

Dr. rer. nat. Jonas Geiping

Full Name	Jonas Alexander Geiping	Email	jonas@tue.ellis.eu
Address	Max-Planck Ring 4 72076 Tübingen Germany	Website	jonasgeiping.github.io
Date of Birth	25th of March 1992, Berlin	Google Scholar	206vNCEAAAAJ

Biography

Jonas Geiping leads a joint research group at the Max Planck Institute for Intelligent Systems and the ELLIS Institute Tübingen. He is a faculty member of the Tübingen AI Center and currently serving on the steering committee of the center. He is a member of the ELLIS society (the European Laboratory for Learning and Intelligent Systems) as an ELLIS scholar.

His research group is investigating questions of safety and efficiency in modern machine learning, with a particular focus on modern language models.

Research

Oct. 2023 ELLIS Institute Tübingen & Max-Planck for Intelligent Systems, Tübingen AI Center
- *Hector Endowed ELLIS Fellow*

Principal investigator at the ELLIS institute through endowment from the Hector II foundation and independent group leader at the Max-Planck for Intelligent Systems through support from the Tübingen AI Center.

Sept. 2021 - University of Maryland, College Park
Oct. 2023 *Postdoctoral Researcher*

Research into privacy and security in machine learning and into deep learning as a science, from an optimization perspective with applications in federated learning, stochastic training of neural networks and topics in security.

May 2021 University of Siegen
July 2021 *Postdoctoral Associate*

Postdoctoral Research: optimization and generalization in deep learning.

Oct. 2016 University of Siegen
March 2021 *Research Associate*

Research in the fields of mathematical optimization, machine learning and computer vision with Prof. Michael Möller with work done in non-convex composite optimization, convex relaxations, optimization for computer vision, learning of optimization objectives and bi-level optimization problems in security.

Aug. 2019 University of Maryland, College Park
Nov. 2019 *Short-Term Scholar Exchange Visitor*

Visiting the group of Prof. Tom Goldstein, joint research in topics of data poisoning for machine learning models and empirical evaluation of deep learning theory.

Oct 2014 - University of Münster (WWU) - Cells-in-Motion Cluster of Excellence
June 2016 Research Assistant
CiM flexible fund project FF-2014-06 - " [Analysis of cell-cell interactions during neuronal migration in the developing cortex by live cell imaging and cell shape quantification](#)"

Education

Dec 2016 - University of Siegen
April 2021 Dr. rer. nat. (PhD), Computer Science
Advisor: Prof. Michael Möller
Thesis: *Modern Optimization Techniques in Computer Vision - From Variational Models to Machine Learning Security*

Oct 2014 - University of Münster
Sept. 2016 M.Sc., Mathematics
Advisor: Prof. Martin Burger
Thesis: *Image Analysis of Neural Tissue Development: Variational Methods for Segmentation and 3D-Reconstruction from large pinhole confocal fluorescence microscopy*

Oct 2011 - University of Münster
Sept. 2014 B.Sc., Mathematics
Advisor: Prof. Martin Burger
Thesis: *Topology-Preserving Segmentation Methods and Application to Mitotic Cell Tracking*

Selected Publications

Jonas Geiping and Tom Goldstein. Cramming: Training a Language Model on a single GPU in one day. In *Proceedings of the 40th International Conference on Machine Learning*, pages 11117–11143. PMLR, July 2023. URL <https://proceedings.mlr.press/v202/geiping23a.html>.

Jonas Geiping, Hartmut Bauermeister, Hannah Dröge, and Michael Moeller. Inverting Gradients - How easy is it to break privacy in federated learning? In *Advances in Neural Information Processing Systems*, volume 33, December 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/c4ede56bbd98819ae6112b20ac6bf145-Abstract.html>.

Jonas Geiping, Micah Goldblum, Phil Pope, Michael Moeller, and Tom Goldstein. Stochastic Training is Not Necessary for Generalization. In *International Conference on Learning Representations*, September 2021. URL <https://openreview.net/forum?id=ZBESeIUB5k>.

Jonas Geiping, Alex Stein, Manli Shu, Khalid Saifullah, Yuxin Wen, and Tom Goldstein. Coercing LLMs to do and reveal (almost) anything. In *SeT LLM @ ICLR 2024*, Vienna, Austria, February 2024. URL <http://arxiv.org/abs/2402.14020>.

Jonas Geiping, Sean McLeish, Neel Jain, John Kirchenbauer, Siddharth Singh, Brian R. Bartoldson, Bhavya Kailkhura, Abhinav Bhatle, and Tom Goldstein. Scaling up Test-Time Compute with Latent Reasoning: A Recurrent Depth Approach. *arxiv:2502.05171[cs]*, February 2025. doi: 10.48550/arXiv.2502.05171. URL <http://arxiv.org/abs/2502.05171>.

John Kirchenbauer, **Jonas Geiping**, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. A Watermark for Large Language Models. In *Proceedings of the 40th International Conference on Machine Learning*, pages 17061–17084. PMLR, July 2023. URL <https://proceedings.mlr.press/v202/kirchenbauer23a.html>.

Recently in the News

- April 2025** [Matthew Berman, forwardfuture.ai](#)
New AI Model "Thinks" Without Using a Single Token
- April 2025** [Quanta Magazine](#)
To Make Language Models Work Better, Researchers Sidestep Language
- Feb 2025** [ntv.de](#)
So fährt Chinas Deepseek US-Techriesen in die KI-Parade
- Jan 2025** [Science Media Germany](#)
Skalierung und Reasoning: Wie geht es mit Sprachmodellen weiter?
- Oct 2024** [Die Zeit](#)
KI-generierte Medien: Zaubertinte gegen die KI-Flut
- Oct 2024** [Tagesspiegel](#)
Warum es bei KI-Wasserzeichen nur zögerlich vorangeht
- June 2024** [Neue Zürcher Zeitung](#)
Chat-GPT, bitte schreibe einen Artikel über Korruption in der Ukraine...
- Jan 2024** [Business Insider](#)
A new AI-detection tool may have solved the problem of false positives
- April 2023** [Neue Zürcher Zeitung](#)
Können Sie unterscheiden, ob ein Text von einem Menschen oder von Chat-GPT kommt?
- Feb 2023** [New York Times](#)
How ChatGPT Could Embed a 'Watermark' in the Text It Generates
- Feb 2023** [Computerphile](#)
Ch(e)at GPT? - Computerphile
- Feb 2023** [WIRED](#)
How to Detect AI-Generated Text, According to Researchers
- Feb 2023** [Nature - News Feature](#)
What ChatGPT and generative AI mean for science
- Feb 2023** [heise.de](#)
Wie Wasserzeichen Texte von KI-Chatbots sichtbar machen könnten
- Jan 2023** [Marktechpost](#)
New AI Research from the University of Maryland Investigates Cramming Challenge
- Jan 2023** [MIT Technology Review](#)
A watermark for chatbots can expose text written by an AI

Full List of Publications (Alphabetic Order)

- Niccolò Ajroldi, Antonio Orvieto, and **Jonas Geiping**. When, Where and Why to Average Weights? In *Proceedings of ICML 2025*, Vancouver, Canada, June 2025. PMLR. URL <https://openreview.net/forum?id=JN8001IZYR>.
- Arpit Bansal, Eitan Borgnia, Hong-Min Chu, Jie S. Li, Hamid Kazemi, Furong Huang, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Cold Diffusion: Inverting Arbitrary Image Transforms Without Noise. In *Thirty-Seventh Conference on Neural Information Processing Systems*, November 2023a. URL <https://openreview.net/forum?id=XH3ArcntI>.
- Arpit Bansal, Hong-Min Chu, Avi Schwarzschild, Soumyadip Sengupta, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Universal Guidance for Diffusion Models. In *The Twelfth International Conference on Learning Representations*, October 2023b. URL <https://openreview.net/forum?id=pzpWBbnwiJ>.
- Valentyn Boreiko, Alexander Panfilov, Vaclav Voracek, Matthias Hein, and **Jonas Geiping**. An Interpretable N-gram Perplexity Threat Model for Large Language Model Jailbreaks. In *Proceedings of ICML 2025*, Vancouver, Canada, June 2025. URL <https://openreview.net/forum?id=ckZbP606Bt>.

- Eitan Borgnia, Valeriia Cherepanova, Liam Fowl, Amin Ghiasi, **Jonas Geiping**, Micah Goldblum, Tom Goldstein, and Arjun Gupta. Strong Data Augmentation Sanitizes Poisoning and Backdoor Attacks Without an Accuracy Tradeoff. In *ICASSP 2021 - 2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 3855–3859, June 2021a. doi: 10.1109/ICASSP39728.2021.9414862.
- Eitan Borgnia, **Jonas Geiping**, Valeriia Cherepanova, Liam Fowl, Arjun Gupta, Amin Ghiasi, Furong Huang, Micah Goldblum, and Tom Goldstein. DP-InstaHide: Provably Defusing Poisoning and Backdoor Attacks with Differentially Private Data Augmentations. In *ICLR 2021 Workshop on Security and Safety in Machine Learning Systems*, March 2021b. URL <http://arxiv.org/abs/2103.02079>.
- Valeriia Cherepanova, Roman Levin, Gowthami Somepalli, **Jonas Geiping**, C. Bayan Bruss, Andrew Gordon Wilson, Tom Goldstein, and Micah Goldblum. A Performance-Driven Benchmark for Feature Selection in Tabular Deep Learning. In *Thirty-Seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, November 2023. URL [https://openreview.net/forum?id=v4PMCdSaAT&referrer=%5BAuthor%20Console%5D\(%2Fgroup%3Fid%3DneurIPS.cc%2F2023%2FTrack%2FDatasets_and_Benchmarks%2FAuthors%23your-submissions\)](https://openreview.net/forum?id=v4PMCdSaAT&referrer=%5BAuthor%20Console%5D(%2Fgroup%3Fid%3DneurIPS.cc%2F2023%2FTrack%2FDatasets_and_Benchmarks%2FAuthors%23your-submissions)).
- Ping-Yeh Chiang, **Jonas Geiping**, Micah Goldblum, Tom Goldstein, Renkun Ni, Steven Reich, and Ali Shafahi. Witchcraft: Efficient PGD Attacks with Random Step Size. In *ICASSP 2020 - 2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 3747–3751, May 2020. doi: 10.1109/ICASSP40776.2020.9052930.
- Ping-yeh Chiang, Renkun Ni, David Yu Miller, Arpit Bansal, **Jonas Geiping**, Micah Goldblum, and Tom Goldstein. Loss Landscapes are All You Need: Neural Network Generalization Can Be Explained Without the Implicit Bias of Gradient Descent. In *The Eleventh International Conference on Learning Representations*, February 2023. URL <https://openreview.net/forum?id=QC10RmRbZy9>.
- Hong-Min Chu, **Jonas Geiping**, Liam H. Fowl, Micah Goldblum, and Tom Goldstein. Panning for Gold in Federated Learning: Targeted Text Extraction under Arbitrarily Large-Scale Aggregation. In *International Conference on Learning Representations*, February 2023. URL <https://openreview.net/forum?id=A9WQaxYsfx>.
- Fay Elhassan, Niccolò Ajroldi, Antonio Orvieto, and **Jonas Geiping**. Can you Finetune your Binoculars? Embedding Text Watermarks into the Weights of Large Language Models. In *ICLR 2025 Workshop on GenAI Watermarking (WMARK)*, April 2025. doi: 10.48550/arXiv.2504.06446. URL <https://openreview.net/forum?id=d1v2hkB0Vj>.
- Liam Fowl, Ping-yeh Chiang, Micah Goldblum, **Jonas Geiping**, Arpit Bansal, Wojtek Czaja, and Tom Goldstein. Preventing Unauthorized Use of Proprietary Data: Poisoning for Secure Dataset Release. In *ICLR 2021 Workshop on Security and Safety in Machine Learning Systems*, February 2021a. URL <http://arxiv.org/abs/2103.02683>.
- Liam Fowl, **Jonas Geiping**, Wojciech Czaja, Micah Goldblum, and Tom Goldstein. Robbing the Fed: Directly Obtaining Private Data in Federated Learning with Modified Models. In *International Conference on Learning Representations*, September 2021b. URL <https://openreview.net/forum?id=fwzUgo0FM9v>.
- Liam Fowl, Micah Goldblum, Ping-yeh Chiang, **Jonas Geiping**, Wojciech Czaja, and Tom Goldstein. Adversarial Examples Make Strong Poisons. In *Advances in Neural Information Processing Systems*, volume 34, pages 30339–30351. Curran Associates, Inc., 2021c. URL <https://proceedings.neurips.cc/paper/2021/hash/fe87435d12ef7642af67d9bc82a8b3cd-Abstract.html>.
- Liam H. Fowl, **Jonas Geiping**, Steven Reich, Yuxin Wen, Wojciech Czaja, Micah Goldblum, and Tom Goldstein. Decepticons: Corrupted Transformers Breach Privacy in Federated Learning for Language Models. In *International Conference on Learning Representations*, February 2023. URL <https://openreview.net/forum?id=r0BrY4BiEX0>.
- Kanchana Vaishnavi Gandikota, **Jonas Geiping**, Zorah Löhner, Adam Czapliński, and Michael Moeller. A Simple Strategy to Provable Invariance via Orbit Mapping. In *Asian Conference on Computer Vision (ACCV)*, Macau, December 2022. arXiv. doi: 10.48550/arXiv.2209.11916. URL <http://arxiv.org/abs/2209.11916>.
- Jonas Geiping** and Tom Goldstein. Cramming: Training a Language Model on a single GPU in one day. In *Proceedings of the 40th International Conference on Machine Learning*, pages 11117–11143. PMLR, July 2023. URL <https://proceedings.mlr.press/v202/geiping23a.html>.
- Jonas Geiping** and Michael Moeller. Composite Optimization by Nonconvex Majorization-Minimization. *SIAM Journal on Imaging Sciences*, pages 2494–2528, January 2018. doi: 10.1137/18M1171989. URL <https://epubs.siam.org/doi/10.1137/18M1171989>.
- Jonas Geiping** and Michael Moeller. Parametric Majorization for Data-Driven Energy Minimization Methods. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 10262–10273, 2019. URL http://openaccess.thecvf.com/content_ICCV_2019/html/Geiping_Parametric_Majorization_for_Data-Driven_Energy_Minimization_Methods_ICCV_2019_paper.html.

- Jonas Geiping**, Hendrik Dirks, Daniel Cremers, and Michael Moeller. Multiframe Motion Coupling for Video Super Resolution. In Marcello Pelillo and Edwin Hancock, editors, *Energy Minimization Methods in Computer Vision and Pattern Recognition*, Lecture Notes in Computer Science, pages 123–138. Springer International Publishing, 2018. ISBN 978-3-319-78199-0. doi: 10.1007/978-3-319-78199-0_9.
- Jonas Geiping**, Hartmut Bauermeister, Hannah Dröge, and Michael Moeller. Inverting Gradients - How easy is it to break privacy in federated learning? In *Advances in Neural Information Processing Systems*, volume 33, December 2020a. URL <https://proceedings.neurips.cc/paper/2020/hash/c4ede56bbd98819ae6112b20ac6bf145-Abstract.html>.
- Jonas Geiping**, Fjedor Gaede, Hartmut Bauermeister, and Michael Moeller. Fast Convex Relaxations using Graph Discretizations. In *31st British Machine Vision Conference (BMVC 2020, Oral Presentation)*, Virtual, September 2020b. URL https://www.bmvc2020-conference.com/conference/papers/paper_0694.html.
- Jonas Geiping**, Liam Fowl, Gowthami Somepalli, Micah Goldblum, Michael Moeller, and Tom Goldstein. What Doesn't Kill You Makes You Robust(er): Adversarial Training against Poisons and Backdoors. In *ICLR 2021 Workshop on Security and Safety in Machine Learning Systems*, February 2021a. URL <http://arxiv.org/abs/2102.13624>.
- Jonas Geiping**, Liam H. Fowl, W. Ronny Huang, Wojciech Czaja, Gavin Taylor, Michael Moeller, and Tom Goldstein. Witches' Brew: Industrial Scale Data Poisoning via Gradient Matching. In *International Conference on Learning Representations*, April 2021b. URL <https://openreview.net/forum?id=01olnfLIbD>.
- Jonas Geiping**, Micah Goldblum, Phil Pope, Michael Moeller, and Tom Goldstein. Stochastic Training is Not Necessary for Generalization. In *International Conference on Learning Representations*, September 2021c. URL <https://openreview.net/forum?id=ZBESeIUB5k>.
- Jonas Geiping**, Jovita Lukasik, Margret Keuper, and Michael Moeller. DARTS for Inverse Problems: A Study on Hyperparameter Sensitivity. *arXiv:2108.05647 [cs]*, August 2021d. URL <http://arxiv.org/abs/2108.05647>.
- Jonas Geiping**, Micah Goldblum, Gowthami Somepalli, Ravid Shwartz-Ziv, Tom Goldstein, and Andrew Gordon Wilson. How Much Data Are Augmentations Worth? An Investigation into Scaling Laws, Invariance, and Implicit Regularization. In *International Conference on Learning Representations*, February 2023. URL <https://openreview.net/forum?id=3aQs3MCSEXD>.
- Jonas Geiping**, Alex Stein, Manli Shu, Khalid Saifullah, Yuxin Wen, and Tom Goldstein. Coercing LLMs to do and reveal (almost) anything. In *SeT LLM @ ICLR 2024*, Vienna, Austria, February 2024. URL <http://arxiv.org/abs/2402.14020>.
- Soumya Suvra Ghosal, Souradip Chakraborty, **Jonas Geiping**, Furong Huang, Dinesh Manocha, and Amrit Bedi. A Survey on the Possibilities & Impossibilities of AI-generated Text Detection. *Transactions on Machine Learning Research*, October 2023. ISSN 2835-8856. URL <https://openreview.net/forum?id=AXtFeYjboj#tab-your-consoles>.
- Shashwat Goel, Joschka Strüder, Ilze Amanda Auzina, Karuna K. Chandra, Ponnurangam Kumaraguru, Douwe Kiela, Ameya Prabhu, Matthias Bethge, and **Jonas Geiping**. Great Models Think Alike and this Undermines AI Oversight. In *Proceedings of ICML 2025*, Vancouver, Canada, June 2025. PMLR. URL <https://openreview.net/forum?id=3Z827FtMNe¬eId=6VHfKYl2hg>.
- Micah Goldblum, **Jonas Geiping**, Avi Schwarzschild, Michael Moeller, and Tom Goldstein. Truth or backpropaganda? An empirical investigation of deep learning theory. In *Eighth International Conference on Learning Representations (ICLR 2020, Oral Presentation)*, April 2020. URL https://iclr.cc/virtual_2020/poster_HyxyIgHFvr.html.
- Andreas Görlitz, **Jonas Geiping**, and Andreas Kolb. Piecewise Rigid Scene Flow with Implicit Motion Segmentation. In *2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 1758–1765, November 2019. doi: 10.1109/IROS40897.2019.8968018.
- Abhimanyu Hans, John Kirchenbauer, Yuxin Wen, Neel Jain, Hamid Kazemi, Prajwal Singhanian, Siddharth Singh, Gowthami Somepalli, **Jonas Geiping**, Abhinav Bhatele, and Tom Goldstein. Be like a Goldfish, Don't Memorize! Mitigating Memorization in Generative LLMs. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, September 2024a. URL <https://openreview.net/forum?id=DylSyAfmWs>.
- Abhimanyu Hans, Avi Schwarzschild, Valeriia Cherepanova, Hamid Kazemi, Aniruddha Saha, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Spotting LLMs With Binoculars: Zero-Shot Detection of Machine-Generated Text. In *Proceedings of the Forty-first International Conference on Machine Learning*, January 2024b. URL <https://openreview.net/forum?id=pXzPMrjjqG>.

- W. Ronny Huang, **Jonas Geiping**, Liam Fowl, Gavin Taylor, and Tom Goldstein. MetaPoison: Practical General-purpose Clean-label Data Poisoning. In *Advances in Neural Information Processing Systems*, volume 33, Vancouver, Canada, December 2020. URL https://proceedings.neurips.cc/paper_files/paper/2020/hash/8ce6fc704072e351679ac97d4a985574-Abstract.html.
- Neel Jain, Khalid Saifullah, Yuxin Wen, John Kirchenbauer, Manli Shu, Aniruddha Saha, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Bring Your Own Data! Self-Sensitivity Evaluation for Large Language Models. In *First Conference on Language Modeling*, August 2024. URL <https://openreview.net/forum?id=k2xZYPZo34#discussion>.
- John Kirchenbauer, **Jonas Geiping**, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. A Watermark for Large Language Models. In *Proceedings of the 40th International Conference on Machine Learning*, pages 17061–17084. PMLR, July 2023a. URL <https://proceedings.mlr.press/v202/kirchenbauer23a.html>.
- John Kirchenbauer, **Jonas Geiping**, Yuxin Wen, Manli Shu, Khalid Saifullah, Kezhi Kong, Kasun Fernando, Aniruddha Saha, Micah Goldblum, and Tom Goldstein. On the Reliability of Watermarks for Large Language Models. In *The Twelfth International Conference on Learning Representations*, October 2023b. URL <https://openreview.net/forum?id=DEJIDCmW0z>.
- John Kirchenbauer, Garrett Honke, Gowthami Somepalli, **Jonas Geiping**, Katherine Lee, Daphne Ippolito, Tom Goldstein, and David Andre. LMD3: Language Model Data Density Dependence. In *First Conference on Language Modeling*, August 2024. URL <https://openreview.net/forum?id=eGCw1UV0hk#discussion>.
- Jie Li, Yow-Ting Shiue, Yong-Siang Shih, and **Jonas Geiping**. Augmenters at SemEval-2023 Task 1: Enhancing CLIP in Handling Compositionality and Ambiguity for Zero-Shot Visual WSD through Prompt Augmentation and Text-To-Image Diffusion. In *Proceedings of the The 17th International Workshop on Semantic Evaluation (SemEval-2023)*, pages 44–49, Toronto, Canada, July 2023. Association for Computational Linguistics. URL <https://aclanthology.org/2023.semeval-1.5>.
- Jie S. Li, **Jonas Geiping**, Micah Goldblum, Aniruddha Saha, and Tom Goldstein. LLM-Generated Passphrases That Are Secure and Easy to Remember. In *The 2025 Annual Conference of the Nations of the Americas Chapter of the ACL*, January 2025. URL <https://openreview.net/forum?id=E68v0F9arG>.
- Jovita Lukasik, **Jonas Geiping**, Michael Moeller, and Margret Keuper. Differentiable Architecture Search: A One-Shot Method? In *AutoML Conference 2023*, August 2023. URL <https://openreview.net/forum?id=LV-5kHj-uV5>.
- Sean Michael McLeish, Arpit Bansal, Alex Stein, Neel Jain, John Kirchenbauer, Brian R. Bartoldson, Bhavya Kailkhura, Abhinav Bhatele, **Jonas Geiping**, Avi Schwarzschild, and Tom Goldstein. Transformers Can Do Arithmetic with the Right Embeddings. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, September 2024. URL <https://openreview.net/forum?id=aIyNLWXuD0>.
- Khalid Saifullah, Yuxin Wen, **Jonas Geiping**, Micah Goldblum, and Tom Goldstein. Seeing in Words: Learning to Classify through Language Bottlenecks. In *ICLR TinyPapers*, May 2023. URL https://openreview.net/forum?id=_QreMdMNIz-.
- Pedro Sandoval-Segura, Vasu Singla, Liam Fowl, **Jonas Geiping**, Micah Goldblum, David Jacobs, and Tom Goldstein. Poisons that are learned faster are more effective. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 197–204, June 2022a. doi: 10.1109/CVPRW56347.2022.00033.
- Pedro Sandoval-Segura, Vasu Singla, **Jonas Geiping**, Micah Goldblum, Tom Goldstein, and David W. Jacobs. Autoregressive Perturbations for Data Poisoning. In *Advances in Neural Information Processing Systems*, December 2022b. URL <https://openreview.net/forum?id=1vusesyN7E>.
- Pedro Sandoval-Segura, Vasu Singla, **Jonas Geiping**, Micah Goldblum, and Tom Goldstein. What Can We Learn from Unlearnable Datasets? In *Thirty-Seventh Conference on Neural Information Processing Systems*, November 2023. URL <https://openreview.net/forum?id=yGs9vTRjaE>.
- Agniv Sharma and **Jonas Geiping**. Efficiently Dispatching Flash Attention For Partially Filled Attention Masks. In *ENLSP Workshop at NeurIPS 2024*, September 2024. doi: 10.48550/arXiv.2409.15097. URL <http://arxiv.org/abs/2409.15097>.
- Manli Shu, Jiong Xiao Wang, Chen Zhu, **Jonas Geiping**, Chaowei Xiao, and Tom Goldstein. On the Exploitability of Instruction Tuning. In *Thirty-Seventh Conference on Neural Information Processing Systems*, November 2023. URL <https://openreview.net/forum?id=4AQ4Fnemox>.

- Siddharth Singh, Prajwal Singhania, Aditya Ranjan, John Kirchenbauer, **Jonas Geiping**, Yuxin Wen, Neel Jain, Abhimanyu Hans, Manli Shu, Aditya Tomar, Tom Goldstein, and Abhinav Bhatele. Democratizing AI: Open-source Scalable LLM Training on GPU-based Supercomputers. In *2024 SC24: International Conference for High Performance Computing, Networking, Storage and Analysis SC*, pages 36–49. IEEE Computer Society, November 2024. ISBN 979-8-3503-5291-7. doi: 10.1109/SC41406.2024.00010. URL <https://www.computer.org/csdl/proceedings-article/sc/2024/529100a036/21HUV5yQsyQ>.
- Vasu Singla, Pedro Sandoval-Segura, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. A Simple and Efficient Baseline for Data Attribution on Images. *arxiv:2311.03386[cs]*, November 2023. doi: 10.48550/arXiv.2311.03386. URL <http://arxiv.org/abs/2311.03386>.
- Shiven Sinha, Shashwat Goel, Ponnurangam Kumaraguru, **Jonas Geiping**, Matthias Bethge, and Ameya Prabhu. Can Language Models Falsify? Evaluating Algorithmic Reasoning with Counterexample Creation. In *Second Conference on Language Modeling*, Montreal, Canada, August 2025. colm.cc. URL <https://openreview.net/forum?id=M7cl4Ldw61>.
- Gowthami Somepalli, Vasu Singla, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Diffusion Art or Digital Forgery? Investigating Data Replication in Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6048–6058, 2023a. URL https://openaccess.thecvf.com/content/CVPR2023/html/Somepalli_Diffusion_Art_or_Digital_Forgery_Investigating_Data_Replication_in_Diffusion_CVPR_2023_paper.html.
- Gowthami Somepalli, Vasu Singla, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Understanding and Mitigating Copying in Diffusion Models. In *Thirty-Seventh Conference on Neural Information Processing Systems*, November 2023b. URL <https://openreview.net/forum?id=HtMXRGbUMt>.
- Gowthami Somepalli, Arkabandhu Chowdhury, **Jonas Geiping**, Ronen Basri, Tom Goldstein, and David W. Jacobs. CALVIN: Improved Contextual Video Captioning via Instruction Tuning. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, September 2024a. URL <https://openreview.net/forum?id=7Kz7icCZ6H>.
- Gowthami Somepalli, Anubhav Gupta, Kamal Gupta, Shramay Palta, Micah Goldblum, **Jonas Geiping**, Abhinav Shrivastava, and Tom Goldstein. Investigating Style Similarity in Diffusion Models. In *Proceedings of the European Conference on Computer Vision*, Milan, April 2024b. arXiv. doi: 10.48550/arXiv.2404.01292. URL <http://arxiv.org/abs/2404.01292>.
- Yuxin Wen, **Jonas Geiping**, Liam Fowl, Micah Goldblum, and Tom Goldstein. Fishing for User Data in Large-Batch Federated Learning via Gradient Magnification. In *Proceedings of the 39th International Conference on Machine Learning*, pages 23668–23684. PMLR, June 2022a. URL <https://proceedings.mlr.press/v162/wen22a.html>.
- Yuxin Wen, **Jonas Geiping**, Liam Fowl, Hossein Souri, Rama Chellappa, Micah Goldblum, and Tom Goldstein. Thinking Two Moves Ahead: Anticipating Other Users Improves Backdoor Attacks in Federated Learning. In *AdvML Frontiers Workshop at 39th International Conference on Machine Learning*, Baltimore, Maryland, USA, 2022b. arXiv. doi: 10.48550/arXiv.2210.09305. URL https://advml-frontier.github.io/pdf/31/CameraReady/backdoor_fl.pdf.
- Yuxin Wen, Arpit Bansal, Hamid Kazemi, Eitan Borgnia, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Canary in a Coalmine: Better Membership Inference with Ensembled Adversarial Queries. In *International Conference on Learning Representations*, February 2023a. URL <https://openreview.net/forum?id=b7SBTEBFnC>.
- Yuxin Wen, **Jonas Geiping**, Micah Goldblum, and Tom Goldstein. STYX: Adaptive Poisoning Attacks Against Byzantine-Robust Defenses in Federated Learning. In *ICASSP 2023 - 2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5, June 2023b. doi: 10.1109/ICASSP49357.2023.10096606.
- Yuxin Wen, Neel Jain, John Kirchenbauer, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. Hard Prompts Made Easy: Gradient-Based Discrete Optimization for Prompt Tuning and Discovery. In *Thirty-Seventh Conference on Neural Information Processing Systems*, November 2023c. URL <https://openreview.net/forum?id=V0stHxDdsN>.
- Yuxin Wen, John Kirchenbauer, **Jonas Geiping**, and Tom Goldstein. Tree-Rings Watermarks: Invisible Fingerprints for Diffusion Images. In *Thirty-Seventh Conference on Neural Information Processing Systems*, November 2023d. URL <https://openreview.net/forum?id=Z57JrmubN1>.
- Yuxin Wen, Leo Marchyok, Sanghyun Hong, **Jonas Geiping**, Tom Goldstein, and Nicholas Carlini. Privacy Backdoors: Enhancing Membership Inference through Poisoning Pre-trained Models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, September 2024. URL <https://openreview.net/forum?id=KppBAWJbry>.

Kaiyu Yue, Bor-Chun Chen, **Jonas Geiping**, Hengduo Li, Tom Goldstein, and Ser-Nam Lim. Object Recognition as Next Token Prediction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 16645–16656, 2024. URL https://openaccess.thecvf.com/content/CVPR2024/html/Yue_Object_Recognition_as_Next-Token_Prediction_CVPR_2024_paper.html.

Recent Preprints

Randall Balestriero, Mark Ibrahim, Vlad Sobal, Ari Morcos, Shashank Shekhar, Tom Goldstein, Florian Bordes, Adrien Bardes, Gregoire Mialon, Yuandong Tian, Avi Schwarzschild, Andrew Gordon Wilson, **Jonas Geiping**, Quentin Garrido, Pierre Fernandez, Amir Bar, Hamed Pirsiavash, Yann LeCun, and Micah Goldblum. A Cookbook of Self-Supervised Learning. *arxiv:2304.12210[cs]*, April 2023. doi: 10.48550/arXiv.2304.12210. URL <http://arxiv.org/abs/2304.12210>.

Nikhil Chandak, Shashwat Goel, Ameya Prabhu, Moritz Hardt, and **Jonas Geiping**. Answer Matching Outperforms Multiple Choice for Language Model Evaluation. *arxiv:2507.02856[cs]*, July 2025. doi: 10.48550/arXiv.2507.02856. URL <http://arxiv.org/abs/2507.02856>.

Jonas Geiping, Sean McLeish, Neel Jain, John Kirchenbauer, Siddharth Singh, Brian R. Bartoldson, Bhavya Kailkhura, Abhinav Bhattele, and Tom Goldstein. Scaling up Test-Time Compute with Latent Reasoning: A Recurrent Depth Approach. *arxiv:2502.05171[cs]*, February 2025. doi: 10.48550/arXiv.2502.05171. URL <http://arxiv.org/abs/2502.05171>.

Neel Jain, Avi Schwarzschild, Yuxin Wen, Gowthami Somepalli, John Kirchenbauer, Ping-yeh Chiang, Micah Goldblum, Aniruddha Saha, **Jonas Geiping**, and Tom Goldstein. Baseline Defenses for Adversarial Attacks Against Aligned Language Models. *arxiv:2309.00614[cs]*, September 2023. doi: 10.48550/arXiv.2309.00614. URL <http://arxiv.org/abs/2309.00614>.

Alexander Panfilov, Paul Kassianik, Maksym Andriushchenko, and **Jonas Geiping**. Capability-Based Scaling Laws for LLM Red-Teaming. *arxiv:2505.20162[cs]*, May 2025. doi: 10.48550/arXiv.2505.20162. URL <http://arxiv.org/abs/2505.20162>.

Xiangyu Qi, Yangsibo Huang, Yi Zeng, Edoardo DeBenedetti, **Jonas Geiping**, Luxi He, Kaixuan Huang, Udari Madhushani, Vikash Sehwal, Weijia Shi, Boyi Wei, Tinghao Xie, Danqi Chen, Pin-Yu Chen, Jeffrey Ding, Ruoxi Jia, Jiaqi Ma, Arvind Narayanan, Weijie J. Su, Mengdi Wang, Chaowei Xiao, Bo Li, Dawn Song, Peter Henderson, and Prateek Mittal. AI Risk Management Should Incorporate Both Safety and Security. *arxiv:2405.19524[cs]*, May 2024. URL <http://arxiv.org/abs/2405.19524>.

Vasu Singla, Pedro Sandoval-Segura, Micah Goldblum, **Jonas Geiping**, and Tom Goldstein. A Simple and Efficient Baseline for Data Attribution on Images. *arxiv:2311.03386[cs]*, November 2023. doi: 10.48550/arXiv.2311.03386. URL <http://arxiv.org/abs/2311.03386>.

Hossein Souri, Arpit Bansal, Hamid Kazemi, Liam Fowl, Aniruddha Saha, **Jonas Geiping**, Andrew Gordon Wilson, Rama Chellappa, Tom Goldstein, and Micah Goldblum. Generating Potent Poisons and Backdoors from Scratch with Guided Diffusion. *arxiv:2403.16365[cs]*, March 2024. URL <http://arxiv.org/abs/2403.16365>.

Teodora Srećković, **Jonas Geiping**, and Antonio Orvieto. Is your batch size the problem? Revisiting the Adam-SGD gap in language modeling. *arxiv:2506.12543[cs]*, June 2025. doi: 10.48550/arXiv.2506.12543. URL <http://arxiv.org/abs/2506.12543>.

Guinan Su and **Jonas Geiping**. Fine, I'll Merge It Myself: A Multi-Fidelity Framework for Automated Model Merging. *arxiv:2502.04030[cs]*, February 2025. doi: 10.48550/arXiv.2502.04030. URL <http://arxiv.org/abs/2502.04030>.

Guinan Su, Li Shen, Lu Yin, Shiwei Liu, Yanwu Yang, and **Jonas Geiping**. GPTailor: Large Language Model Pruning Through Layer Cutting and Stitching. *arxiv:2506.20480[cs]*, June 2025. doi: 10.48550/arXiv.2506.20480. URL <http://arxiv.org/abs/2506.20480>.

Yanwu Yang, Guinan Su, Jiesi Hu, Francesco Sammarco, **Jonas Geiping**, and Thomas Wolfers. MedSAMix: A Training-Free Model Merging Approach for Medical Image Segmentation. *arxiv:2508.11032[cs]*, August 2025. doi: 10.48550/arXiv.2508.11032. URL <http://arxiv.org/abs/2508.11032>.